

CS3243 Tetris Project

Varun Gupta Gao Bo Advay Pal Chang Chu-Ming Herbert Ilhan Tanujaya
Varun Gupta Gao Bo Advay Pal Chang Chu-Ming Herbert Ilhan Tanujaya

1 Introduction

Tetris is likely one of the world's most famous and popular games. In this report, we describe how we devise an agent to play the game of Tetris. We use an agent that greedily picks the best possible next state from a given state, by using a heuristic function to approximate the value of a state. To train our heuristic function, we use a novel algorithm that is a combination of the well known Genetic Algorithm (GA) and Particle Swarm Algorithm (PSO). Our agent manages to clear **TODO** million lines on average with a max of **TODO** million lines, demonstrating that our algorithm is effective.

2 Agent Strategy

Our agent uses a linear weighted sum of features as the heuristic function for a given state. Here, a state is defined by the configuration of the board. Given a state and a piece, the agent computes the heuristic value for all possible next states, and then greedily picks the next state with the maximum heuristic value. In other words, the next state is chosen by

$$\max_{s' \in N(s,p)} \left\{ \sum w_i f_i(s') \right\},$$

where s is the original state, $N(s,p)$ is the set of all possible next states from s and a piece p , $f_i(s')$ is the score of feature i on state s' and w_i is the weight of feature i .

3 Features

We used the following features for our heuristic function:

- **Altitude Difference** ($AltD$): Difference between height of the highest and the lowest column
- **Column Transition** ($ColT$): Number of adjacent squares in a column with opposite parity (where parity is defined as either full or empty)
- **Deepest Well** (DW): Height of the lowest column
- **Height** ($ColH$): Height of the highest column
- **Number of Columns With Holes** (CwH): The number of columns with holes, where a hole is defined as an empty square directly beneath a filled square
- **Number of Holes** ($NHole$): Number of holes in the entire board
- **Number of Wells** ($NWell$): Number of columns that have height less than the 2 adjacent columns
- **Rows Cleared** (RC): Number of rows cleared for that particular move
- **Row Transition** ($RowT$): Number of adjacent squares in a row with opposite parity
- **Total Column Height** ($TColH$): The sum of the heights of all the columns

- **Total Column Height Difference** (*TColHD*): The sum of the difference of heights between adjacent columns
- **Weighted Block** (*WB*): Sum of value of every filled square, where a square's value is equal to the row it is in (numbered from the bottom)
- **Well Sum** (*Wells*): Sums up the depth of every well

4 Our Algorithm

Our algorithm consists of a combination of a GA and PSO algorithm. We run both of these algorithms on different islands with population sizes of a 100 each. Each member of our population is a heuristic (set of weights). The key idea is that both of the algorithms run individually, but after each generation, we copy the 10 best heuristics on each island to the other one.

GA

We use a standard GA with crossover and mutation. A crossover is done by taking two parent heuristics, and, for every feature, selecting the weight from one of the parents randomly. A mutation is defined as multiplying the current weight of a feature with a value normally distributed with a mean of 1 and a standard deviation of 1. Parents are chosen based on fitness, with fitter population members having a higher probability of being chosen. We keep the top half of the population to ensure that each set of heuristics that perform well will remain within the population and replace the bottom half with new members.

PSO algorithm

Again, we use a standard PSO algorithm which treats every heuristic as a point in 13-dimensional space. This is the position of the heuristic. Each heuristic also has a velocity, a vector in the 13-dimensional space, which is added to the position to get new population members. New velocities are computed by taking a linear weighted sum of the old velocity, the personal best of that member, and the global best.

Rationale behind combination of algorithms

We first ran PSO and GA individually. In doing so, we realised that PSO was in general not doing very well, but that it sometimes made a leap and jumped by almost an order of magnitude in terms of lines cleared. On the other hand, GA seemed to keep getting better, but only quite slowly. Thus we hypothesised that GA may be doing more of exploitation, finding the maximum in the local area, while PSO may have been doing more of exploration, looking for good solutions all over the search space. Hence we thought that perhaps if we combined the two, PSO could lead GA to better areas, which GA could then drill down into. This would allow us to achieve a more balanced trade-off between exploitation and exploration, which we hoped would make our algorithm more efficient.

5 Further Optimisation

The hyperparameters for both the GA and the PSO algorithm need to be optimized for our integrated solution to learn at an optimal rate. Given limited time, we decided to focus our attention on the hyperparameters of the PSO algorithm. The rationale behind this is that, since the PSO algorithm looks to be exploring the state space, if we could successfully optimize the PSO algorithm, it would be able to find a solution close to the global maximum and share this information with the GA. The GA would

then refine the solution further.

The hyperparameters of the PSO algorithm consist of ρ_g and ρ_p , the constants which determine the effect of the global and personal bests on each iteration of velocity change; and ω , the constant of inertia. Typically **TODO**, ρ_g and ρ_p are both set to 2.0, which leaves ω to be optimized. A high ω will favour exploration, while a low ω will favour exploitation. Based on the observations of [Shi and Eberhart, 1998], we decided to first test values of ω between 0.6 to 1.2. While initializations with low ω generally performed better, the results were not conclusive. Thus, we decided that instead of a fixed ω , we could vary it based on how close the population is to the global maximum. This is estimated by taking \log_{10} of the average fitness scores of the population. The formula is as follows:

$$\omega = 1.2 - k(\log_{10}(\text{Average Fitness}) - 3)$$

We tested different values of k by running multiple jobs for an hour and comparing the average of the fitness scores. Through experimentation, we found that the algorithm performs best when k is 0.2.

6 Experiments and Analysis

6.1 Learning Performance

We ran our learning algorithm on nodes provided by the National Supercomputing Centre cluster. The nodes we used had the following architecture specifications: E5-2690 v3 @ 2.60GHz. Using this, the results we obtained were:

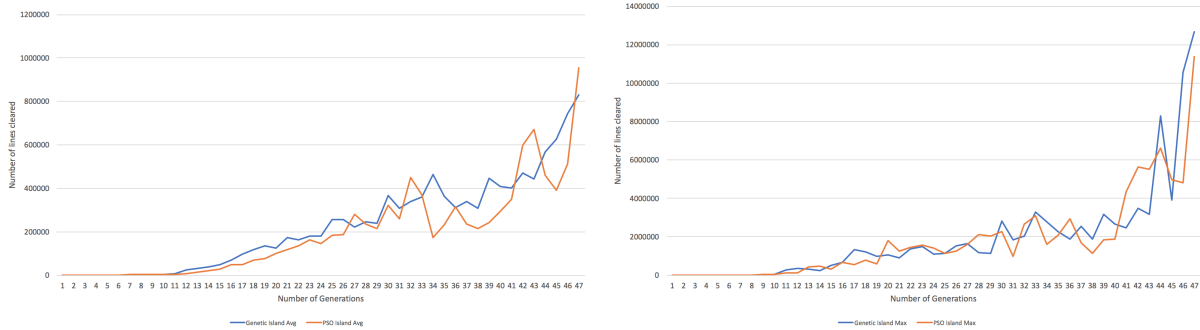


Figure 1: Performance of the Learning Algorithm

The first figure plots the lines cleared by the best heuristic in the GA and PSO islands over 47 generations, while the second figure plots the average lines cleared by the same populations. It can be seen that sometimes PSO leads the way in the first figure, producing a good heuristic that GA then builds upon. However, this is not always true, and in fact we did observe in other runs of the algorithm that sometimes GA also led the way.

We also compared the results of our algorithm with Naive GA and Naive PSO (i.e. we ran them in isolation). The results obtained were:

This figure plots the maximum of the average of GA and PSO islands for our learning algorithm, and compares it with the averages of Naive GA and Naive PSO. It can be seen that our algorithm performed far better, with PSO getting stuck at around 20 thousand lines cleared and GA rising very slowly.

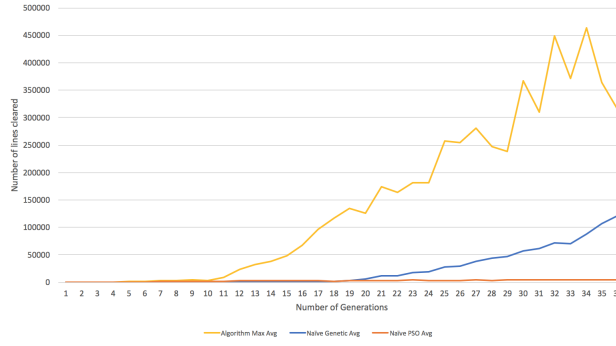


Figure 2: Comparison of Learning Algorithm with Naive GA/PSO

However, we can also observe from Figure 1 that there is a large amount of variability in the results of the algorithm. One reason for this is the element of randomness in the algorithm itself, for example, in the mutation amount and threshold in GA and in the initial velocity imparted to a particle in PSO. There is a large amount of randomness in the game of tetris itself, as seen more clearly in the section analysing our agent’s performance.

Using a version of our algorithm that used 2 threads (one per island), we were able to get to a maximum of around 4 million lines cleared in around a day. Running a parallelised version (see below) we got to a maximum of around 18m in 10 hours.

6.2 Agent Performance

The best weights we obtained were:

<i>AltD</i>	<i>ColT</i>	<i>DW</i>	<i>ColH</i>	<i>CwH</i>	<i>NHole</i>	<i>NWell</i>	<i>RC</i>	<i>RowT</i>	<i>TColH</i>	<i>TColHD</i>	<i>WB</i>	<i>WellS</i>
0	0	0	0	0	0	0	0	0	0	0	0	0

We ran the heuristic on **TODO** games, encapsulated in the following graph:

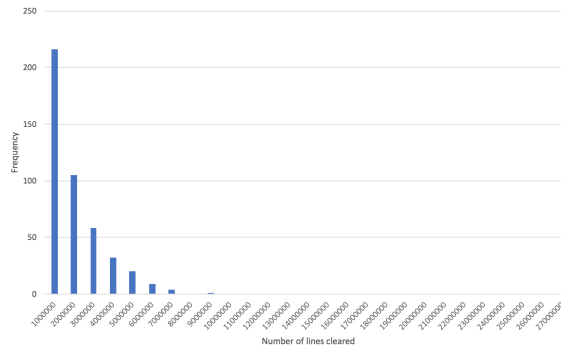


Figure 3: Performance of our Agent

Maximum	Minimum	Average
0	0	0

As can be seen in Figure 3, there is a large variability in how well the heuristic performs, hitting a maximum of **TODO** and a minimum of **TODO**. This is likely because of the large randomisation within the game of tetris, with the sequences of pieces having a large impact on the performance. For example, a long sequence of *S* and *Z* pieces is impossible to get through successfully.

7 Scaling to Big Data

In order to scale to Big Data, we considered how our algorithm performed in the presence of multiple cores. We found that parallelising our algorithms through the use of multithreading lead to a linear speedup, demonstrating that our algorithm scaled very well in the presence of multiple cores. We attribute this to the fact that both GA and PSO algorithms are embarassingly parallel, which is why we were able to achieve an optimal speedup quite easily.

We parallelised our algorithm by running the PSO and GA islands on different cores, and further parallelised it by calculating the fitness of each population member on different cores. We did this as the major computational cost the algorithm was incurring came from the calculation of the fitness values, since that involved running entire tetris games. We started both algorithms as two separate jobs on single machines in the NSCC cluster (which provide 12 cores each) at the same time, and allowed them to run for around 20 hours before comparing the results. We used the total number of lines cleared in the entire time span as a rough estimate of the total amount of CPU time used. Our results showed that the single-threaded algorithm cleared a total of 158182993 lines, while the multi-threaded version cleared a total of 1899941699 lines. This shows that parallelisation provides a speedup of 12 times, which is optimal since each machine has 12 cores.

8 Conclusion

9 References

[Shi and Eberhart, 1998] Shi, Y. and Eberhart, R. (1998). Parameter selection in particle swarm optimization. In *Evolutionary programming VII*, pages 591–600. Springer.